Weather Temperature Forecasting on Jena Climate data

- Time series

Group Members:

Dinesh Chandra Gaddam

Project Overview

Scope of the Project

- Develop the best possible model using Arma, ARIMA, SARIMA or Box-Jenkins to forecast Temperature

- Perform all the required analysis on pretrained and post-trained data

Understanding the concepts and underlying methodology of the data

DATA SOURCE

Public Datasets: Use dataset from Kaggle named Jena Climate Data. Data used of period 2009 to 2012.

Project content

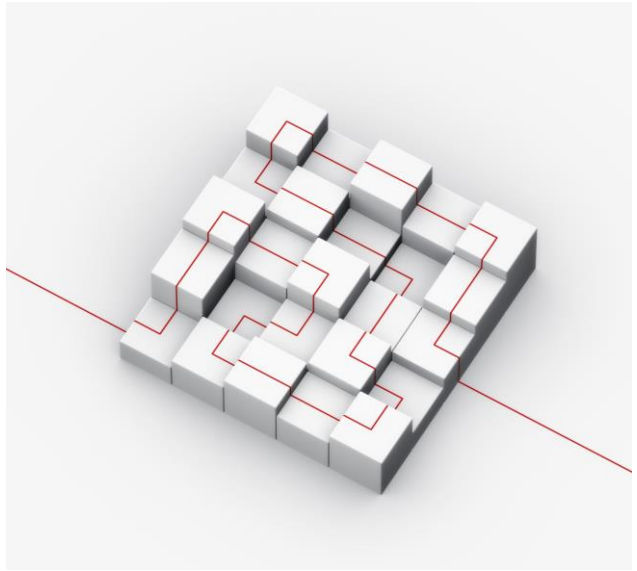

- Data Collection and Preprocessing
- Stationarity Check and Seasonal Decomposition
- Feature Selection and Base Models
- Model Development
- Model Evaluation
- Forecasting
- Deciding best model
- Summary

Data Collection and Preprocessing

I have Limited Data to 3 years from 2009 to 2011 to make the data flexible to work with

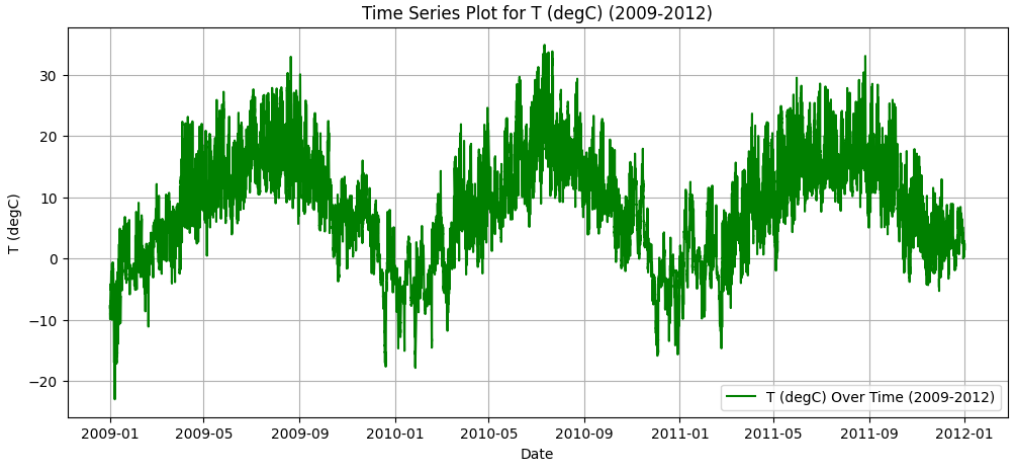

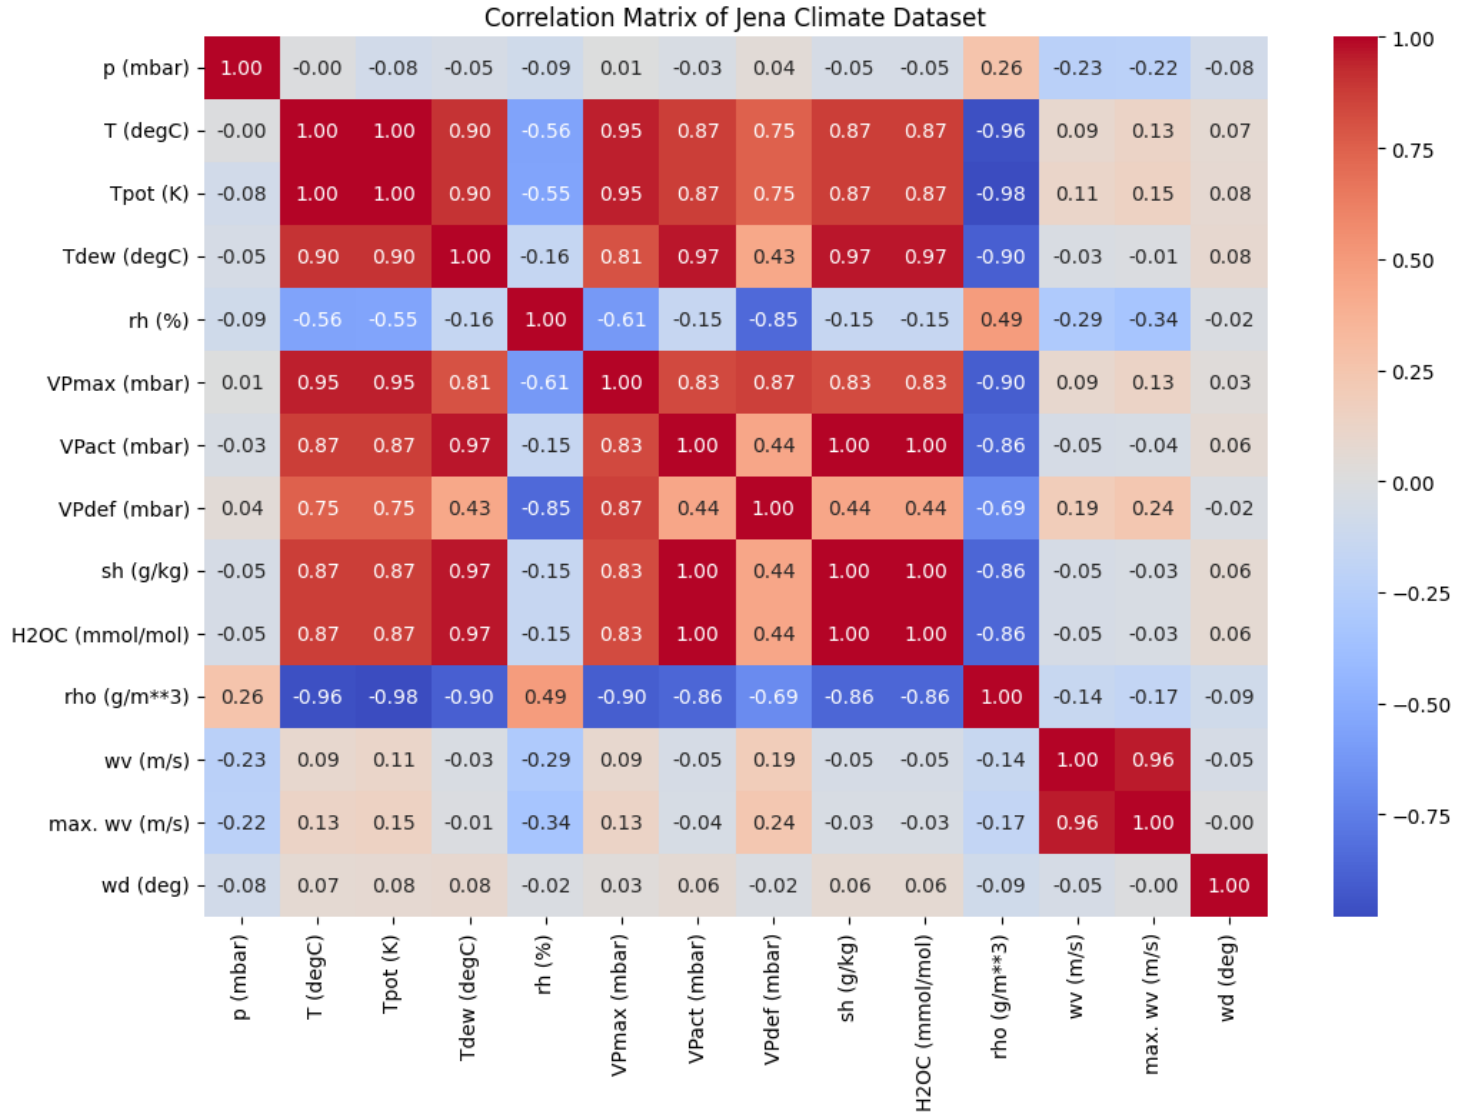

Stationarity Check and Seasonal Decomposition

Before differencing

```
=== Augmented Dickey-Fuller (ADF) Test ===
ADF Statistic: -8.2050
p-value: 0.0000
Critical Values:
  1%: -3.4304
  5%: -2.8616
  10%: -2.5668
✅ The series is likely stationary (reject H0).

=== Kwiatkowski-Phillips-Schmidt-Shin (KPSS) Test ===
KPSS Statistic: 1.9481
p-value: 0.0100
Critical Values:
  10%: 0.3470
  5%: 0.4630
  2.5%: 0.5740
  1%: 0.7390
❌ The series is likely non-stationary (reject H0).
```

After Differencing

```
=== Augmented Dickey-Fuller (ADF) Test ===
ADF Statistic: -76.4568
p-value: 0.0000
Critical Values:
  1%: -3.4304
  5%: -2.8616
  10%: -2.5668
✅ The series is likely stationary (reject H0).

=== Kwiatkowski-Phillips-Schmidt-Shin (KPSS) Test ===
KPSS Statistic: 0.0068
p-value: 0.1000
Critical Values:
  10%: 0.3470
  5%: 0.4630
  2.5%: 0.5740
  1%: 0.7390
✅ The series is likely stationary (fail to reject H0).
```

The Data is Seasonal: with 87% seasonality

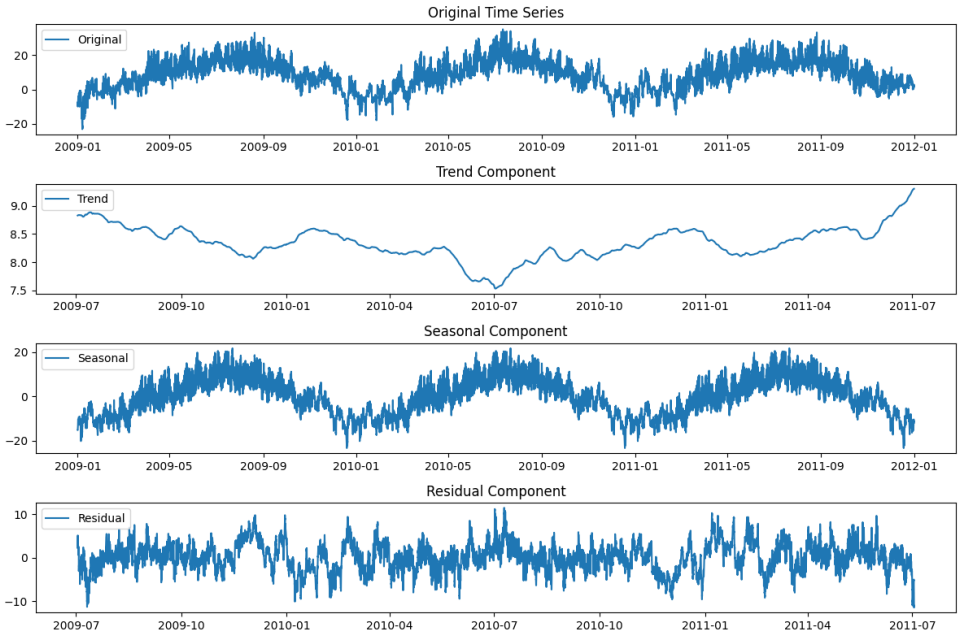

Feature Selection and Base Models

```
Final VIF Results:
         Feature       VIF
4      wd (deg)  4.722527
1        rh (%)  4.439649
0   Tdew (degC)  1.978726
2  VPdef (mbar)  1.948563
3       wv (m/s)  1.000941
```

```
=== FEATURE SELECTION REPORT ===

1. VIF Analysis Results:
   - Removed 6 features due to high VIF
   - Highest remaining VIF: 4.72

2. PCA/SVD Findings:
   - Condition number: 4.55
   - PCA reduced to 4 components (95% variance)

3. Backward Stepwise Regression:
   - Selected 7 features
   - Final features: ['Tdew (degC)', 'rh (%)', 'VPdef (mbar)', 'wv (m/s)', 'wd (deg)', 'hour', 'day_of_year']
```

Feature Selection and Base Models

Model Comparison:

	Model	MSE	RMSE
3	SES	25.164543	5.016427
1	Naïve	25.164815	5.016454
2	Drift	25.321412	5.032039
0	Average	55.381009	7.441842

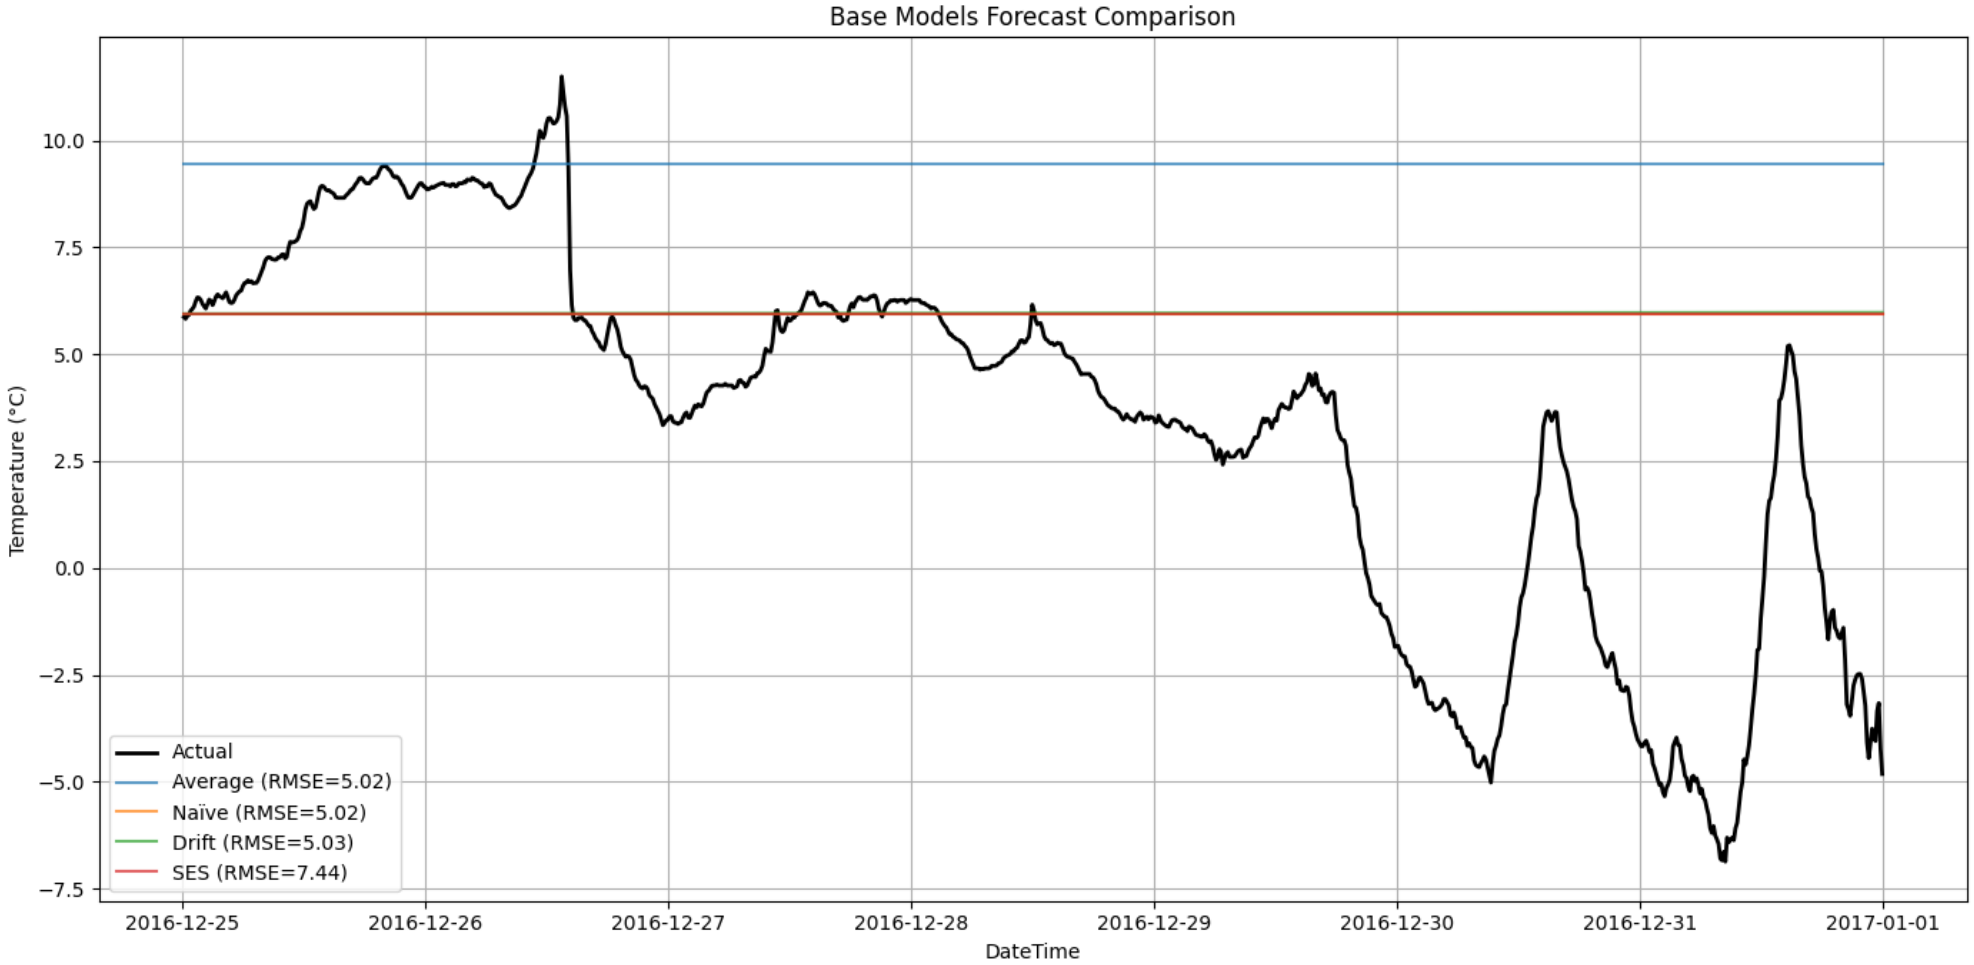

Model Development

Multiple Linear Regression

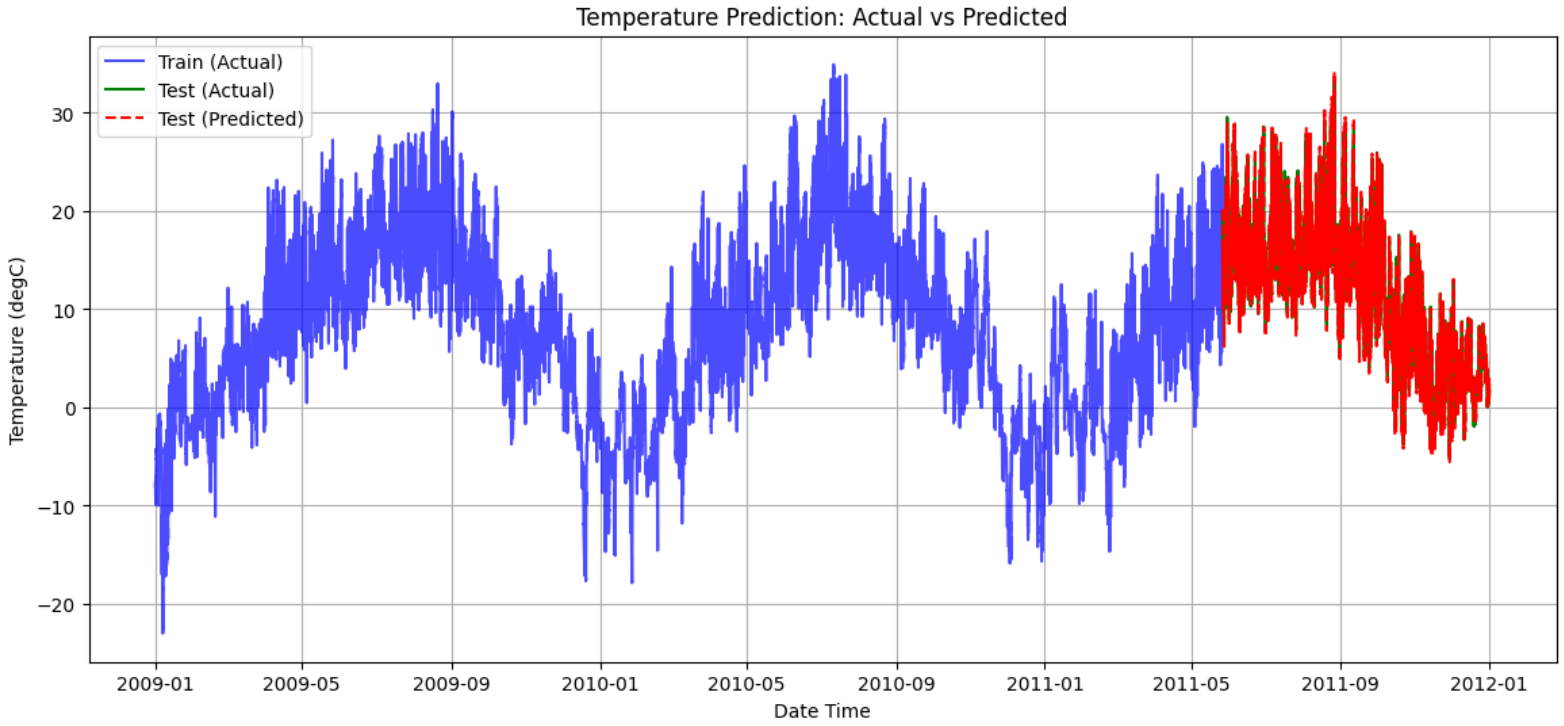


```
=== CROSS VALIDATION RESULTS ===
Mean MSE: 0.2934
Std MSE: 0.1870
```

```
=== COMPLETE REGRESSION ANALYSIS ===
                            OLS Regression Results
==============================================================================
Dep. Variable:                T (degC)   R-squared:                       0.997
Model:                             OLS   Adj. R-squared:                  0.997
Method:                  Least Squares   F-statistic:                 7.115e+06
Date:                 Sat, 03 May 2025   Prob (F-statistic):               0.00
Time:                         00:00:13   Log-Likelihood:                -77024.
No. Observations:               126256   AIC:                         1.541e+05
Df Residuals:                   126248   BIC:                         1.541e+05
Df Model:                            7
Covariance Type:             nonrobust
==============================================================================
                 coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
const         16.5825      0.016   1068.444      0.000      16.552      16.613
Tdew (degC)    0.9802      0.000   2871.759      0.000       0.979       0.981
rh (%)        -0.1705      0.000  -1009.613      0.000      -0.171      -0.170
VPdef (mbar)   0.2561      0.001    411.722      0.000       0.255       0.257
wv (m/s)      -0.0083      0.001     -9.426      0.000      -0.010      -0.007
hour          -0.0027      0.000    -13.843      0.000      -0.003      -0.002
day_sin        0.0466      0.002     22.950      0.000       0.043       0.051
day_cos       -0.0132      0.003     -3.905      0.000      -0.020      -0.007
==============================================================================
Omnibus:                    99423.645   Durbin-Watson:                   0.012
Prob(Omnibus):                  0.000   Jarque-Bera (JB):          6054765.930
Skew:                           3.319   Prob(JB):                         0.00
Kurtosis:                      36.270   Cond. No.                         981.
==============================================================================
```

Model Development

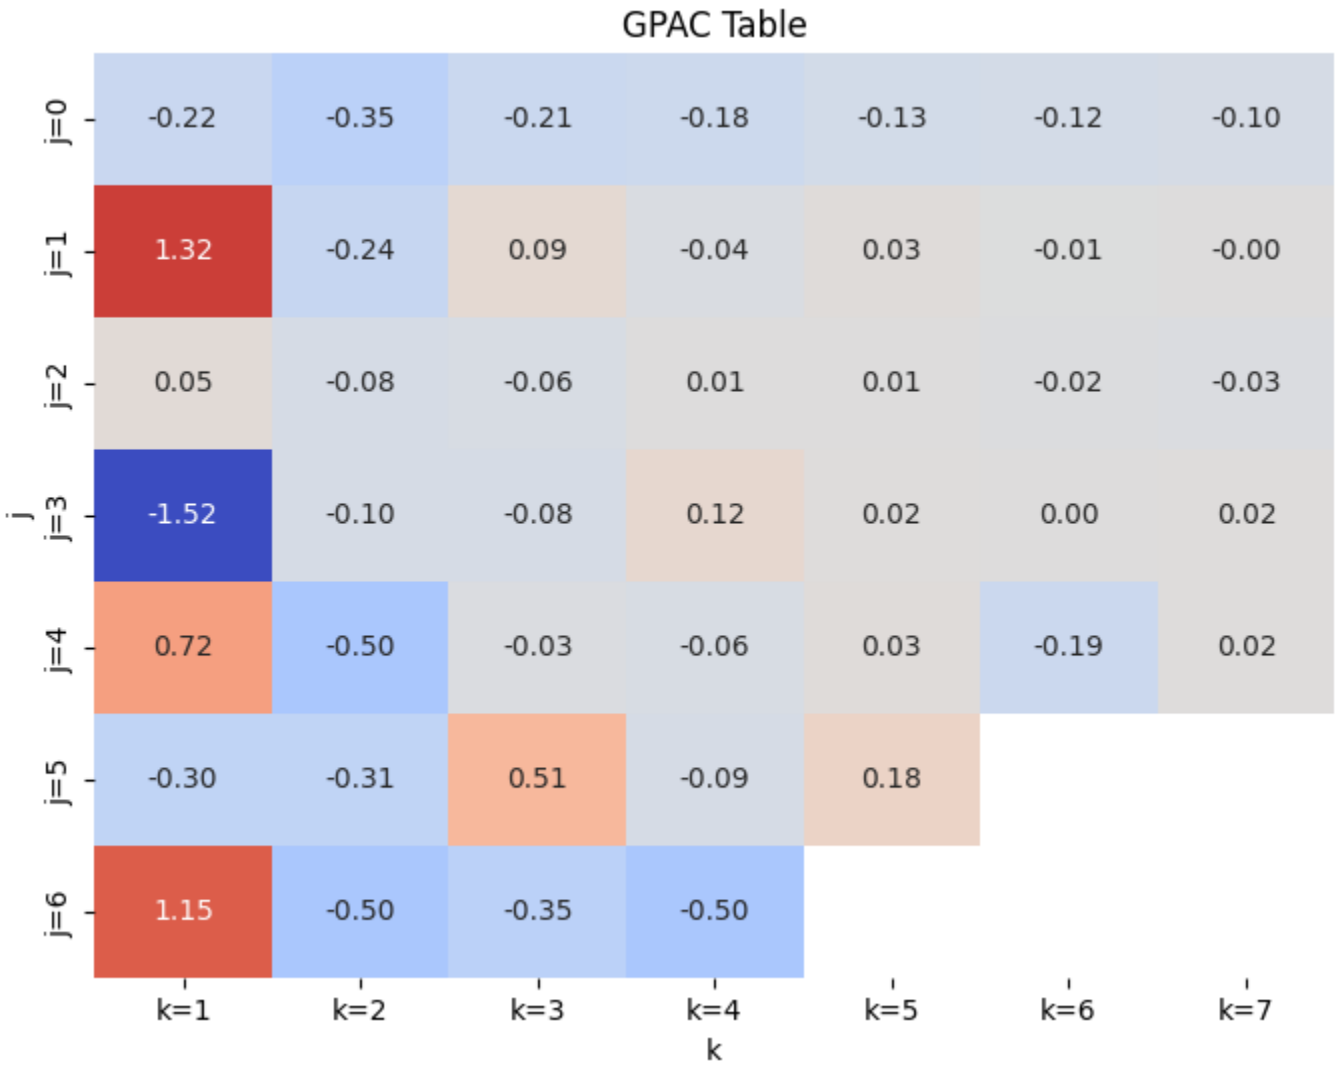

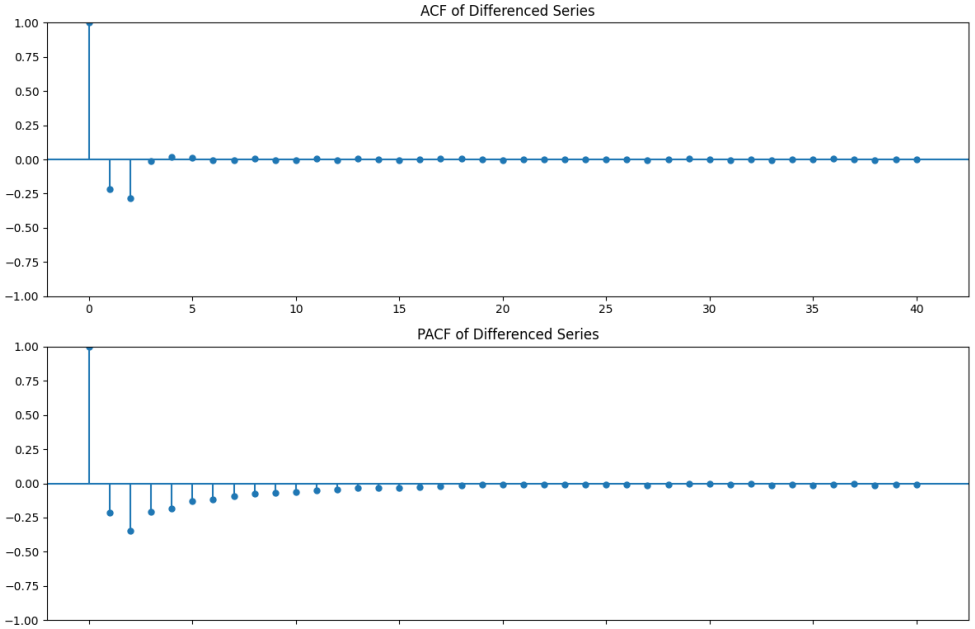


```
[-0.04339786 -0.435999   -0.38974652]
[-0.47316144 -0.36253648]
```

GPAC, PACF and ACF of the Stationary data, expecting order of AR: 0 or 2 and MA order of 2

ARMA observations

```
ARMA(1,2) Q-test:

--- Q-Test Summary ---
Q-statistic                : 1594245.9841
Chi-square Critical (α=0.05, dof=47) : 64.0011
Result                     : ❌ Residuals show autocorrelation (Q > Q*)
```

```
ARMA(0,2) Q-test:

--- Q-Test Summary ---
Q-statistic                : 1617700.4048
Chi-square Critical (α=0.05, dof=48) : 65.1708
Result                     : ❌ Residuals show autocorrelation (Q > Q*)
```

ARIMA Observations

```
==============================================================================
Dep. Variable:                T (degC)   No. Observations:               126256
Model:                 ARIMA(1, 1, 2)   Log Likelihood                20496.334
Date:                Sun, 04 May 2025   AIC                            -40984.667
Time:                        13:52:48   BIC                            -40945.683
Sample:                             0   HQIC                           -40972.959
                               - 126256
Covariance Type:                  opg
==============================================================================
                 coef    std err          z      P>|z|      [0.025      0.975]
------------------------------------------------------------------------------
ar.L1          0.9551      0.001    830.387      0.000       0.953       0.957
ma.L1         -0.4482      0.002   -250.833      0.000      -0.452      -0.445
ma.L2         -0.3389      0.002   -191.935      0.000      -0.342      -0.335
sigma2         0.0423   6.19e-05    683.382      0.000       0.042       0.042
==============================================================================
Ljung-Box (L1) (Q):                  14.74   Jarque-Bera (JB):            933729.75
Prob(Q):                              0.00   Prob(JB):                         0.00
Heteroskedasticity (H):               0.75   Skew:                            -0.58
Prob(H) (two-sided):                  0.00   Kurtosis:                        16.27
==============================================================================
```

```
==============================================================================
Dep. Variable:                T (degC)   No. Observations:               126256
Model:                 ARIMA(0, 1, 2)   Log Likelihood                15134.858
Date:                Sun, 04 May 2025   AIC                            -30263.716
Time:                        13:52:59   BIC                            -30234.478
Sample:                             0   HQIC                           -30254.935
                               - 126256
Covariance Type:                  opg
==============================================================================
                 coef    std err          z      P>|z|      [0.025      0.975]
------------------------------------------------------------------------------
ma.L1          0.5590      0.001    405.975      0.000       0.556       0.562
ma.L2          0.1331      0.002     88.655      0.000       0.130       0.136
sigma2         0.0461   7.24e-05    636.139      0.000       0.046       0.046
==============================================================================
Ljung-Box (L1) (Q):                  28.63   Jarque-Bera (JB):            636405.48
Prob(Q):                              0.00   Prob(JB):                         0.00
Heteroskedasticity (H):               0.77   Skew:                            -0.16
Prob(H) (two-sided):                  0.00   Kurtosis:                        13.99
==============================================================================
```

```
--- Q-Test Summary (lags=50, df=47) ---
Q-statistic                : 835.5226
Chi-square Critical (α=0.05, dof=47) : 64.0011
Result                     : ❌ Residuals show autocorrelation (Q > Q*)
```

```
--- Q-Test Summary ---
Q-statistic                : 28737.0758
Chi-square Critical (α=0.05, dof=48) : 65.1708
Result                     : ❌ Residuals show autocorrelation (Q > Q*)
```

SARIMA observations for its best model..

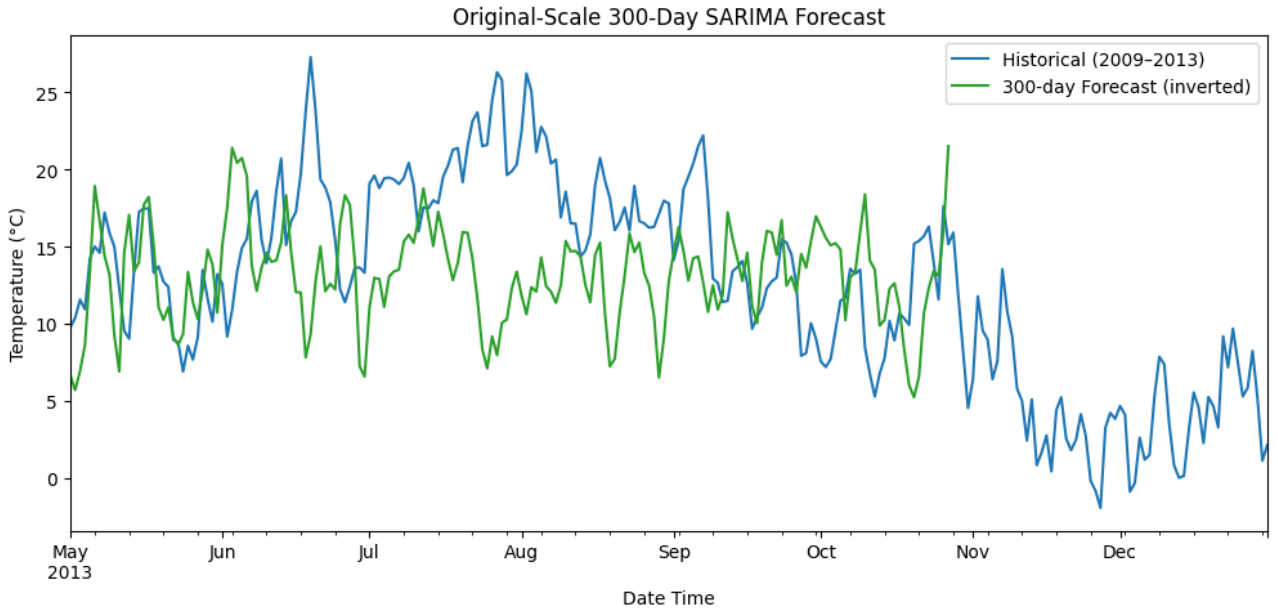

p. Variable:	D.DS365.T (degC)	No. Observations:	1095
del:	SARIMAX(1, 0, 2)x(1, 0, [1], 365)	Log Likelihood	-1797.345
te:	Sun, 04 May 2025	AIC	3606.690
me:	14:27:54	BIC	3634.224
mple:	01-02-2010 - 12-31-2012	HQIC	3617.315
variance Type:	opg		

	coef	std err	z	P>|z|	[0.025	0.975]
.L1	0.7242	0.030	24.246	0.000	0.666	0.783
.L1	-0.7077	14.113	-0.050	0.960	-28.370	26.954
.L2	-0.2923	4.130	-0.071	0.944	-8.388	7.803
.S.L365	-0.4503	0.040	-11.375	0.000	-0.528	-0.373
.S.L365	-0.2885	0.075	-3.866	0.000	-0.435	-0.142
gma2	7.8032	110.167	0.071	0.944	-208.119	223.726

ung-Box (L1) (Q):	0.09	Jarque-Bera (JB):	3.81
ob(Q):	0.76	Prob(JB):	0.15
teroskedasticity (H):	0.88	Skew:	-0.02
ob(H) (two-sided):	0.31	Kurtosis:	3.35

Observation: the model captures the seasonality but there is almost 0 trend in my data. So the model was unable to access the right movement. To solve this issue I used **ARIMA(0,1,2) on Seasonally-Adjusted Data + Seasonal Recomposition Forecast**

ARIMA(0,1,2) on Seasonally-Adjusted Data + Seasonal Recomposition Forecast

Dep. Variable:	y	No. Observations:	1461
Model:	ARIMA(1, 1, 2)	Log Likelihood	-3233.952
Date:	Sun, 04 May 2025	AIC	6475.904
Time:	15:31:21	BIC	6497.049
Sample:	01-01-2009	HQIC	6483.792
	- 12-31-2012		
Covariance Type:	opg		

	coef	std err	z	P>|z|	[0.025	0.975]
ar.L1	0.7445	0.019	39.452	0.000	0.708	0.781
ma.L1	-0.7450	0.027	-27.147	0.000	-0.799	-0.691
ma.L2	-0.2516	0.027	-9.295	0.000	-0.305	-0.199
sigma2	4.9034	0.168	29.205	0.000	4.574	5.233

Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	19.51
Prob(Q):	0.99	Prob(JB):	0.00
Heteroskedasticity (H):	0.92	Skew:	-0.19
Prob(H) (two-sided):	0.35	Kurtosis:	3.43

```
--- Q-Test (lags=50, df=47, alpha=0.05) ---
Q-statistic: 66016.8018
Critical value: 64.0011
Result: ❌ Residuals show autocorrelation
```

365 step forecast of the resampled data=>

```
Forecast metrics: MAE = 3.182 °C, RMSE = 3.999 °C
```

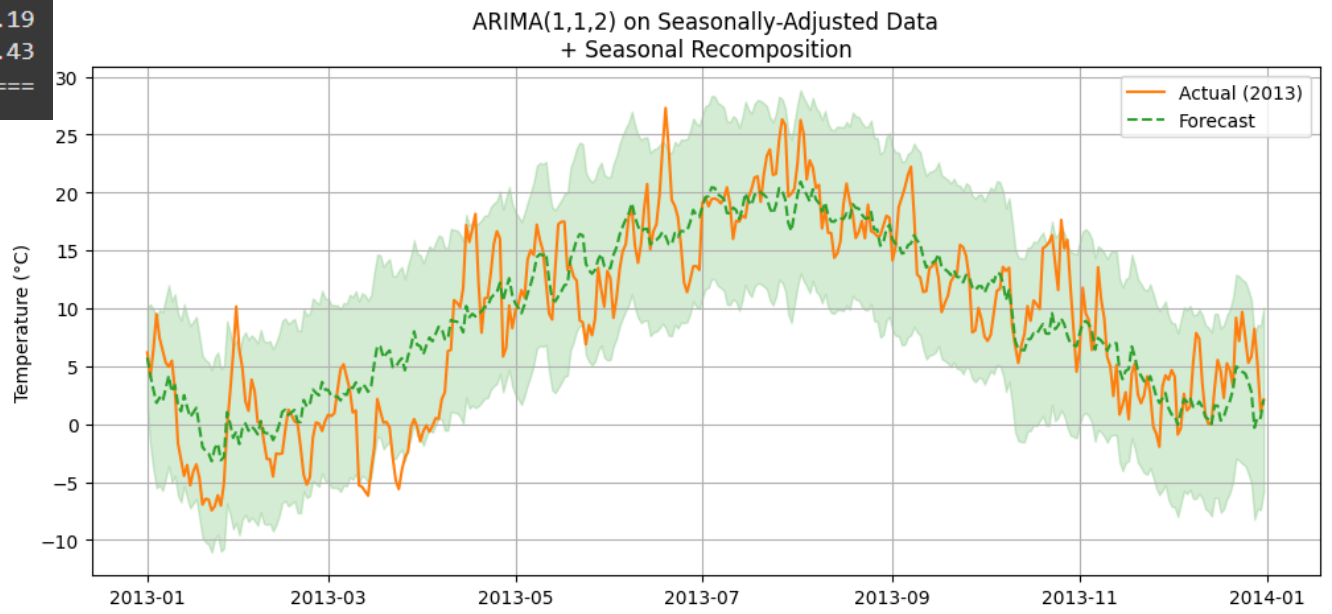

Box-Jenkins

	nb	nf	nc	nd	Q_stat	Q_crit	Q_pass	S_stat	S_crit	S_pass
0	1	1	0	0	143.449470	65.170769	False	21.057283	30.143527	True
1	1	1	0	1	138.611276	64.001112	False	20.187493	30.143527	True
2	1	1	1	0	138.611276	64.001112	False	20.187493	30.143527	True
3	1	1	1	1	138.611276	62.829620	False	20.187493	30.143527	True
4	1	2	0	0	109.112704	64.001112	False	16.472130	28.869299	True
5	1	2	0	1	114.149448	62.829620	False	15.683024	28.869299	True
6	1	2	1	0	114.149448	62.829620	False	15.683024	28.869299	True
7	1	2	1	1	114.149448	61.656233	False	15.683024	28.869299	True
8	2	1	0	0	143.011171	64.001112	False	19.926954	30.143527	True
9	2	1	0	1	137.706429	62.829620	False	20.564540	30.143527	True
10	2	1	1	0	137.706429	62.829620	False	20.564540	30.143527	True
11	2	1	1	1	137.706429	61.656233	False	20.564540	30.143527	True
12	2	2	0	0	109.126378	62.829620	False	15.558826	28.869299	True
13	2	2	0	1	113.870886	61.656233	False	16.051770	28.869299	True
14	2	2	1	0	113.870886	61.656233	False	16.051770	28.869299	True
15	2	2	1	1	113.870886	60.480887	False	16.051770	28.869299	True

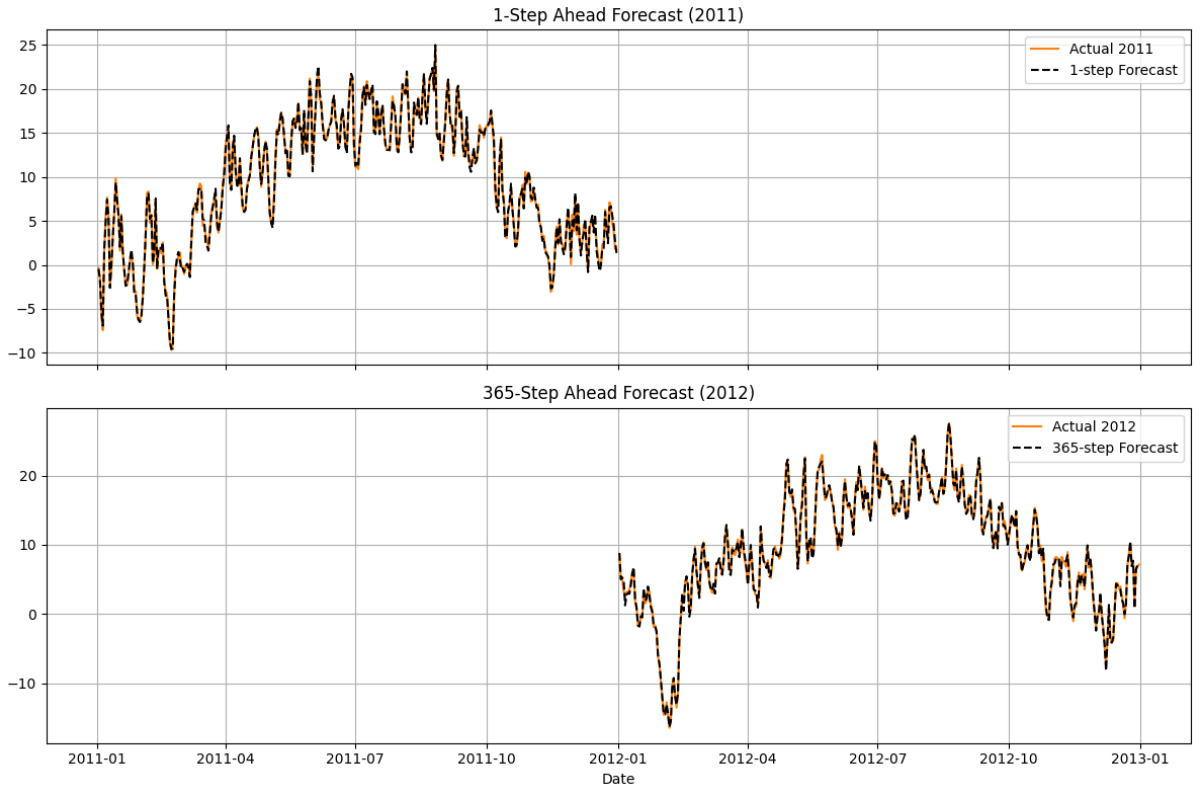

Q-test: Q=143.0, crit=64.0, df=47 -> ❌ autocorrelation

S-test: S=19.9, crit=30.1, df=19 -> ✅ G(q) accurate

Summary

- The Models have been built for ARMA, ARIMA, SARIMA, BOX-Jenkins etc:-al from all of them Box-Jenkins Has performed the best..

- I have derived the order from GPAC, PACF and ACF plots by feeding them the stationary data..

- The model received a accuracy of RMSE of 0.360 for the normalised data after making it stationary...

References

Notes and Assignments

THANK YOU!